

Autocategorization

Andrew Carnes¹

¹ *University of Florida*

Abstract

In order to maximize the chance of discovering a certain effect, many scientific analyses extract subsets of data from the main data set that are sensitive to the effect. If the effect exists, the data in these sensitive subsets is expected to differ significantly from the null hypothesis. There are many ways to extract subsets, but an optimal method should minimize the global p-value on the null hypothesis given the effect. Call the orthogonal subsets of data categories and the selection of subsets categorization. An optimal categorization enhances the power of the experiment and reduces the data needed for discovery. By specifying a metric corresponding to the expected p-value (given the effect), the optimum categorization may be automated using a decision tree. This paper describes the autocategorization algorithm.

1. Introduction

Many scientific analyses attempt to rule out a null hypothesis and declare a discovery. In this context, seeing a large enough discrepancy in the data from the null's prediction leads to a low p-value, the rejection of the null, and a
5 discovery. With a model for the null and another for the expected data, the expected p-value can be quantified. An analysis with a lower p-value represents more convincing evidence to rule out the null, and therefore an analysis with a lower expected p-value is considered more sensitive. In addition, a more sensitive analysis requires less data to rule out the null hypothesis. For both of
10 these reasons, optimizing the sensitivity is critical for any scientist.

The expected hypothesis and the null hypothesis determine probability density functions (PDFs) in a many dimensional feature space. However, data sparsely populates high dimensional spaces, and fits in high dimensions are difficult or untrustworthy in practice, so the PDFs are usually compared over a single discriminating feature - sometimes a few - to determine the p-value. Unfortunately, upon reducing to a lower dimensional PDF comparison, the sensitivity of the analysis is often reduced: feature space regions with large discriminating power are combined with those of high probability but low sensitivity.

To regain sensitivity, the regions of feature space with large discriminating power need to be extracted. In this way, it makes sense to divide the feature space up into separate regions, called categories, and compare the low dimensional PDFs within the categories to maximize the sensitivity. By analytically specifying a metric corresponding to the expected p-value the optimum categorization can be automated. This is separate problem from regression or classification where fit values are applied to different regions of space as to minimize some loss function. The goal is simply to divide up feature space to maximize the statistical sensitivity.

2. The Autocategorizer Algorithm

This section explains the autocategorization algorithm (autocategorizer) in terms of a binned counting experiment using histograms along a single variable, z . Let H_o be the null hypothesis and H_1 be the expected hypothesis – some trusted theory. The goal of the autocategorizer is to partition feature space in order to minimize the expected p-value. To be clear, the expected p-value in this case is the probability to measure something at least as extreme as H_1 given H_o . Minimizing the expected p-value will maximize the chance for discovery if H_1 is true.

Let i label the bin and c label an orthogonal partition of feature space – a category. Each category has its own histograms. Let $z_{c,i}$ describe the amount in the i^{th} bin of the histogram for category c . Moreover, let H_o predict $z_{c,i} = B_{c,i}$

40 in each bin for category c and let H_1 predict $z_{c,i} = S_{c,i} + B_{c,i}$ in each bin for category c .




Figure 1: An example of the S and B histograms in a category. The null hypothesis is given by B , and the expected hypothesis is given by $S+B$.

The likelihood to observe $S_{c,i} + B_{c,i}$ in each bin when $B_{c,i}$ is expected in each bin provides a statistical measure of the discrepancy between H_1 and H_o . A low likelihood corresponds to a large discrepancy and a low p-value. Minimizing the likelihood will minimize the expected p-value. Equivalently, maximizing the negative log-likelihood (NLL) minimizes the expected p-value. The NLL for a single category's histogram is given by,

$$-\log(p(z_{c,i}; \lambda_{c,i})/C_{c,i}) = -\log\left(\prod_i \text{Poisson}(z_{c,i}; \lambda_{c,i})/C_{c,i}\right), \quad (1)$$

where the measured amount in a bin is given by $z_{c,i}$, and the amount expected by the hypothesis is given by $\lambda_{c,i}$.

50 The categories are orthogonal, so the likelihoods for each category multiply, and the net NLL is

$$-\log(p(z_i; \lambda_{c,i})/C_{c,i}) = -\log\left(\prod_{c,i} \text{Poisson}(z_{c,i}; \lambda_{c,i})/C_{c,i}\right). \quad (2)$$

Approximating each Poisson distribution by a Gaussian with $\mu_{c,i} = \lambda_{c,i}$ and $\sigma_{c,i}^2 = \lambda_{c,i}$, leads to a χ^2 variable

$$-\log\left(\prod_i \text{Poisson}(z_{c,i}; \lambda_{c,i})/C_{c,i}\right) = \sum_{c,i} (z_{c,i} - \mu_{c,i})^2 / \sigma_{c,i}^2. \quad (3)$$

The $C_{c,i}$ terms are constants that standardize the χ^2 variable. To calculate the
 55 significance of measuring H_1 given that the null is true, z_i is replaced by $S_i + B_i$
 and μ_i is replaced by B_i . The substitutions provide an approximation for the
 NLL to observe H_1 given H_o ,

$$\text{Net Significance} = \sum_{c,i} (z_{c,i} - \mu_{c,i})^2 / \sigma_{c,i}^2 = \sum_{c,i} S_{c,i}^2 / B_{c,i}. \quad (4)$$

Choosing the categories to maximize the net significance will minimize the p-value.

60 With the net significance acting as the metric, a decision tree[1] can split
 up the feature space as to maximize the metric. The decision tree algorithm
 greedily builds the optimum categorization by recursively splitting feature space
 regions into two using hyperplanes. On the first iteration, the autocategorizer
 calculates the net significance for the inclusive set of all events. The algorithm
 65 then searches over the inclusive set, checking all possible split values of the
 first feature, x_1 . Events with x_1 values less than the split go in one candidate
 category, and those with x_1 values greater than or equal to the split value go into
 the other candidate category. For every split candidate, the algorithm calculates
 the net significance in the two categories delineated by the split value. The x_1
 70 split value that provides the largest gain in significance over the inclusive set
 is stored. The gain is defined in the equation below where c1 and c2 are the
 prospective categories created from c by splitting on the feature.

$$\text{Gain} = \sum_i S_{c1,i}^2 / B_{c1,i} + \sum_i S_{c2,i}^2 / B_{c2,i} - \sum_i S_{c,i}^2 / B_{c,i} \quad (5)$$

The autocategorizer then searches over the second feature, x_2 , and stores the
 split value that provides the largest gain in significance. The process is repeated
 75 for all of the remaining features.

The algorithm chooses to split at the feature value with the largest gain,
 creating two categories from the inclusive set of events. At the next iteration,
 the autocategorizer repeats the procedure for the two new categories and chooses
 to split the category that provides the most gain. This process continues, each

80 time greedily choosing to split the category with the most gain, until the number of categories desired is reached.




Figure 2: An example of the categorization process for features x, y and three categories. The autocategorizer chooses $x=0.52$ for the first split and $y=0.50$ for the second. The colored crosses represent the events that should be grouped together for optimum sensitivity. After three iterations, the categorizer correctly groups those events.

3. Conclusions

The autocategorizer algorithm automatically maximizes the chance to discover an effect of interest in a statistical analysis by extracting sensitive regions from a multidimensional feature space. The algorithm has been used at the Compact Muon Solenoid (CMS) experiment at CERN to set limits on the Higgs particle's rate of decay to two muons [2]. In the Higgs to dimuons analysis, the autocategorizer improves the upper limit on the rate of decay by 15% compared to human expert categorization with the same simulated data and the same features. The improvement is equivalent to collecting 32% more data.

References

1. Breiman L, Friedman J, Olshen R, Stone C. Classification and Regression Trees. Monterey, CA: Wadsworth and Brooks; 1984.
2. CMS Collaboration . Search for the standard model Higgs boson decaying to two muons in pp collisions at $\sqrt{s} = 13$ TeV 2017;URL: <http://cds.cern.ch/record/2292159>.